

AWS Lambda & Serverless Computing 🚀



🚀 What is Serverless?

Serverless doesn't mean there are no servers. It means **you don't manage** the servers! Instead, the cloud provider (like AWS) handles the infrastructure so you can focus on writing code.

- You focus on writing code, while the cloud provider (AWS) handles infrastructure.
- No need to provision, scale, or maintain servers.
- Pay-per-use billing (you pay only for execution time).

☁️ Serverless in AWS (Popular Services)

Here are the main AWS services that support serverless architecture:

AWS Service	What It Does
🧠 AWS Lambda	Run code without managing servers
🌐 Amazon API Gateway	Create & manage RESTful APIs for serverless apps
📦 Amazon DynamoDB	NoSQL database that scales automatically
📦 AWS Fargate	Run containers without managing EC2
📁 Amazon S3	Serverless object storage
🔄 Amazon EventBridge / CloudWatch Events	Trigger functions on events
✉️ Amazon SNS / SQS	Serverless messaging and notifications
🔑 AWS Cognito	Serverless user authentication and access control





AWS Lambda Overview

What is AWS Lambda?

AWS Lambda is a **serverless compute service**. It lets you run your code **without managing servers**. You just upload your code, and AWS runs it automatically **in response to events** — like an API call, file upload to S3, or a change in DynamoDB.

You only pay when your code is running — by **milliseconds**, not hours!

Benefits of AWS Lambda

-  **No server management** – You don't need to launch or maintain EC2 instances or patch operating systems.
-  **Automatic scaling** – Lambda can handle one request or a million without you configuring anything.
-  **Event-driven** – Lambda can automatically trigger on events from S3, API Gateway, DynamoDB, CloudWatch, and more.
-  **Fast and easy to deploy** – Just write your code, upload, and it's ready to run.

♦ Language Support

AWS Lambda supports:

- Node.js, Python, Java, C#, Go, Ruby, PowerShell
- Custom Runtimes (Bring your own language via Docker)

Why AWS Lambda Instead of EC2?

Using EC2 means you manage the **whole virtual server**. You install the OS, handle patches, manage auto scaling, and monitor everything. You're charged **per second or hour**, even when your app is idle.

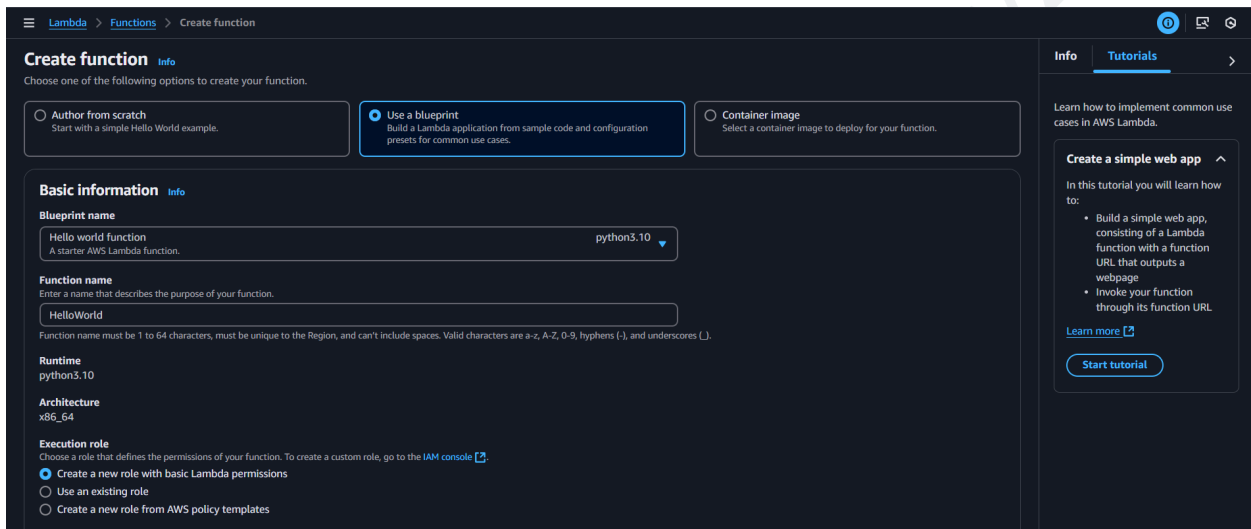
Lambda is different. You only write the **code**, and AWS runs it **automatically** when triggered. There's **no server to manage**, no need to worry about scaling, and you only **pay when your code runs**.

So, if you just need to run short tasks like handling an image upload, processing form data, or reacting to an S3 event — **Lambda is perfect**.

AWS Lambda Hands-On Guide

✓ Step 1: Create the Function

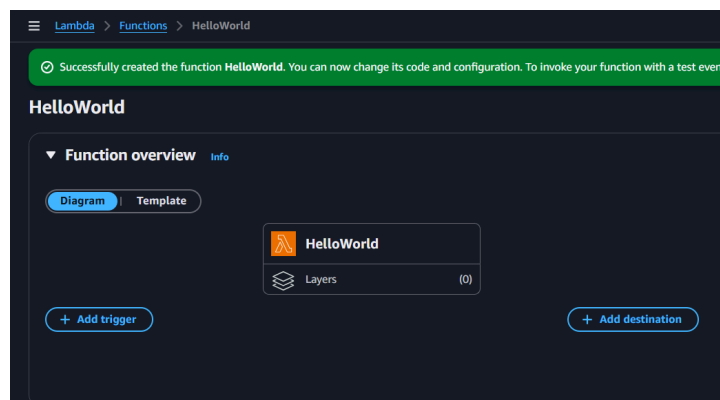
1. Go to **AWS Lambda** service.
2. Click **Create function**.
3. Select **Use a blueprint**.
4. Pick a **Blueprint name** that matches your use case (like `hello-world` or `s3-get-object`).
5. Give your function a **name**.
6. Click **Create function**.

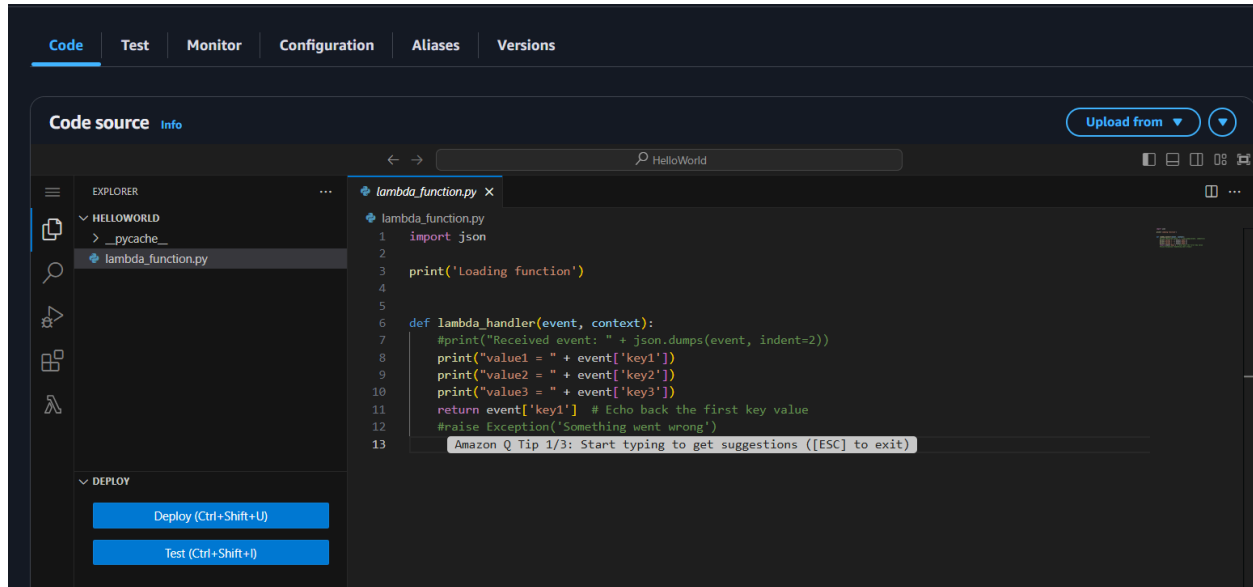


The screenshot shows the AWS Lambda 'Create function' console. The 'Use a blueprint' option is selected. The 'Blueprint name' is 'Hello world function' and the 'Runtime' is 'python3.10'. The 'Function name' is 'HelloWorld'. The 'Execution role' is 'Create a new role with basic Lambda permissions'.

Step 2: View Your Function

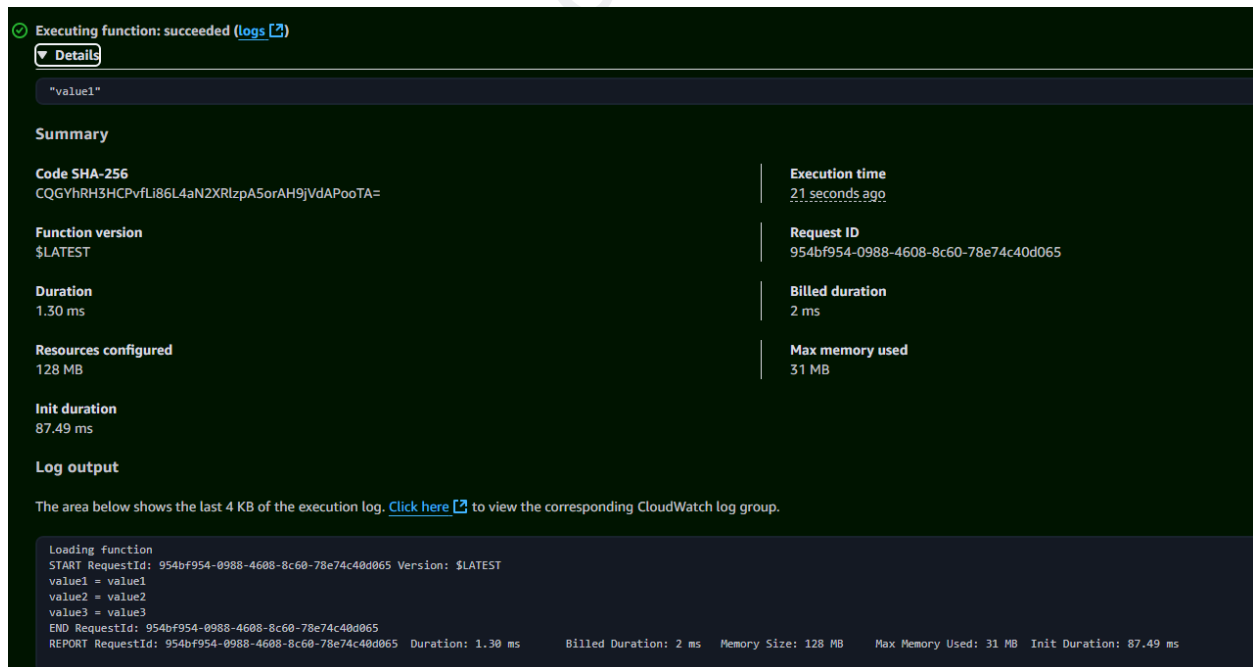
- After creation, you'll see:
 - **Function overview**
 - **Code source editor** to view/edit code.





Step 3: Test the Function

1. Click on the **Test** button.
2. Give your test event a **name**.
3. Leave the default test JSON (or customize it if needed).
4. Click **Save**.
5. Then click **Test** again to run it.





Step 4: Monitor Logs

- Go to the **Monitoring** tab.
- Click **View logs in CloudWatch** to see:
 - Execution details
 - Logs and errors

The screenshot shows the AWS CloudWatch 'Log events' page. The breadcrumb trail is: CloudWatch > Log groups > /aws/lambda/HelloWorld > 2025/05/27/[[\$LATEST]]a20a802dc9144b07a50a8283d62b192d. The page has a search bar with the placeholder 'Filter events - press enter to search'. On the right, there are buttons for 'Actions', 'Start tailing', and 'Create metric filter'. Below the search bar, there are filters for 'Clear', '1m', '30m', '1h', '12h', 'Custom', 'Local timezone', and 'Display'. The log events table has two columns: 'Timestamp' and 'Message'. The messages include: 'No older events at this moment. Retry', 'INIT_START Runtime Version: python:3.10.v76 Runtime Version ARN: arn:aws:lambda:ap-south-1:runtime:3795f2e5ed4539bb3e8e88e9958f8851955bb8e88a1cc2580a5a5dc87e5b5ac', 'Loading function', 'START RequestId: 954bf954-0988-4608-8c60-78e74c40d065 Version: \$LATEST', 'value1 = value1', 'value2 = value2', 'value3 = value3', 'END RequestId: 954bf954-0988-4608-8c60-78e74c40d065', and 'REPORT RequestId: 954bf954-0988-4608-8c60-78e74c40d065 Duration: 1.30 ms Billed Duration: 2 ms Memory Size: 128 MB Max Memory Used: 31 MB Init Duration: 87.49 ms'. At the bottom, it says 'No newer events at this moment. Auto retry paused. Resume'.



Execution Limits

Item	Default Limit
Concurrent executions	1,000 (can be increased)
Memory allocation	128 MB to 10,240 MB (10 GB)
Maximum execution timeout	15 minutes
Environment variables size	4 KB
Ephemeral storage (/tmp)	512 MB (can go up to 10 GB)
Payload size (Request/Response)	6 MB (for synchronous)
Event size (Async)	256 KB



Deployment Limits

Item	Default Limit
Deployment package size (compressed .zip/.jar)	50 MB (direct upload)
Uncompressed deployment size	250 MB
Container image size (for Lambda using ECR)	10 GB
Layers per function	5 layers (max total size: 250 MB unzipped)

What is Concurrency?

Concurrency means:

💡 *How many instances of a function can run at the same time.*

Example:

If 5 users call your Lambda function at the same time 🕒, and your concurrency is 5, AWS will run 5 **parallel** instances of the function — one for each user ✅.

Lambda Concurrency

In AWS Lambda, **Concurrency** =

Number of Lambda executions running at the same time in a region.

Types of Concurrency:

Type	Description
🟢 Unreserved concurrency	Shared across all functions in the region (default: 1,000).
🛡️ Reserved concurrency	You can reserve a fixed number for a specific function.
🔒 Provisioned concurrency	Pre-warms your function (always ready = no cold starts). Great for low-latency needs.

What is Throttle?

Throttle means:

💥 *Blocking or rejecting extra requests when limits are hit.*

Lambda Throttling

Lambda **throttles** function invocations when:

1. 😵 **Concurrent executions exceed the limit** (e.g., 1,000).
2. 🔒 You **reserved concurrency too low** and it's fully used.
3. 📦 **Provisioned concurrency is full** and new requests arrive.

What happens when throttled?

- **Synchronous calls** → Get a **429 - TooManyRequestsException** ❌
- **Asynchronous calls** → Automatically retried **2 times** with delays (if retry is enabled) 🔄

What's a Cold Start?

When Lambda needs to spin up a new instance of your function for the **first time**, it:

1. **Downloads your code**
2. **Initializes dependencies**
3. Then runs your code.

This takes time — especially for Java. That delay is called a **cold start**.



AWS Lambda SnapStart

What is SnapStart?

SnapStart is an AWS Lambda feature that  **takes a snapshot of your function after it's initialized**, so it can **start much faster** later!

It reduces **cold start time** ⌚, especially for **Java** functions — which are known for slow startup.

How SnapStart Works:

1.  During **deployment**, AWS runs your function **once**, initializes it, and **takes a snapshot**.
2.  When a new instance is needed (cold start), Lambda just **restores from the snapshot** — super fast!

What is Customization at the Edge?

Customization at the Edge means running code **closer to users**, in **AWS Edge Locations** around the world — instead of in centralized servers or regions.

This helps with:

- ⚡ Faster response times
- 🛡️ Early request filtering or redirects
- 🛠️ Customizing content per user/location/device

AWS offers two main services for this:

- ✨ **CloudFront Functions**
- 🧠 **Lambda@Edge**

⚡ CloudFront Functions

CloudFront Functions are **lightweight JavaScript functions** that run **at the edge**, right inside Amazon CloudFront.

✅ Use Cases:

- URL rewrites or redirects
- Header manipulation
- Access control (block bots, check geolocation)
- Simple A/B testing

💡 Key Features:

- Super **fast** (runs in microseconds)
- **Low-cost** and **highly scalable**
- Runs **only at the viewer request/response** stage

Lambda@Edge

Lambda@Edge lets you run **AWS Lambda functions** at CloudFront edge locations.

Use Cases:

- Dynamic content generation at the edge
- Authentication & authorization
- Device-based content delivery (e.g., mobile vs desktop)
- Modify requests/responses deeply (e.g., cookies, user-agent logic)

Key Features:

- Supports **Node.js** and **Python**
- Can run at **4 stages**:
 1. Viewer Request
 2. Origin Request
 3. Origin Response
 4. Viewer Response
- Can access the **full request and response**
- More **powerful** than CloudFront Functions, but with **slightly higher latency**




Lambda in VPC

By default, **AWS Lambda runs in its own secure environment** — it doesn't have access to your **VPC resources** like RDS databases, EC2 instances, or private subnets.

But you can configure Lambda to **connect to your VPC**, allowing it to access internal resources (like private DBs) securely.

Why Put Lambda in a VPC?

You need it when your Lambda function needs to:

-  Connect to **private subnets**
-  Access **RDS, ElastiCache**, or **EC2** in a private subnet
-  Use custom **security groups** or **network ACLs**

How It Works:

When you configure Lambda to run inside a VPC, it:

1. Uses **Elastic Network Interfaces (ENIs)** to connect to your VPC.
2. Attaches to the **subnets** and **security groups** you specify.
3. Can access private resources — **but loses default internet access** unless you configure a **NAT Gateway** or **NAT instance**!

AWS Certification Exam Tips:

- Lambda can run in a VPC to access **private resources**.
- Once in a VPC, it **loses default internet access** — unless you configure a **NAT Gateway**.
- Lambda uses **ENIs** to connect to your VPC.
- You must provide **subnet IDs** and **security group IDs** during configuration.